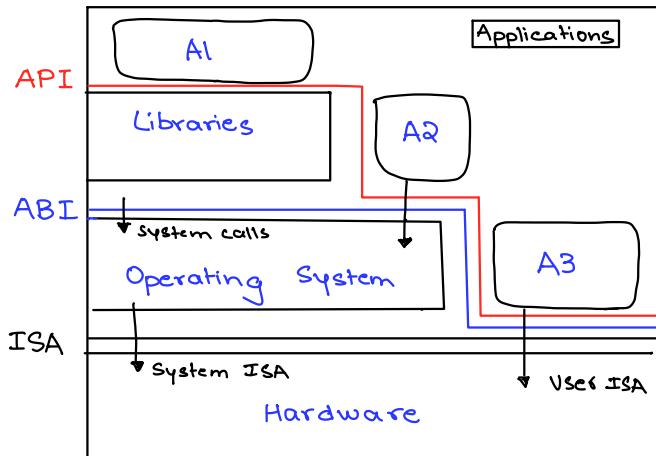




UNIT 3

LAYERING

- managing system complexity - layering
- minimizes interactions
- 2 modes → user & privileged



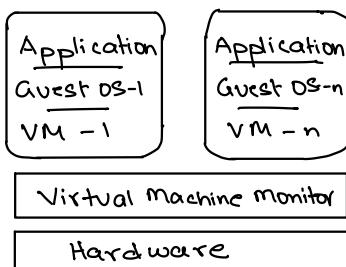
- 1st interface - ISA - boundary b/w s/w & h/w - processor's set of instructions
- ABI (Appⁿ Binary Int.) - H/w access to the appⁿ's & library modules. Doesn't include privileged instructions.
- API - defines set of instr the h/w was designed to execute - gives appⁿ access to the ISA.

CONDITIONS FOR EFFICIENT VIRTUALIZATION

- A program running under the VMM must exhibit behaviour essentially similar to that demonstrated when the appⁿ runs directly on an equivalent machine.
- The VMM should be in complete control of the virtualized resources.
- A statistically significant fraction of the machine instructions should be executed without the intervention of the VMM.

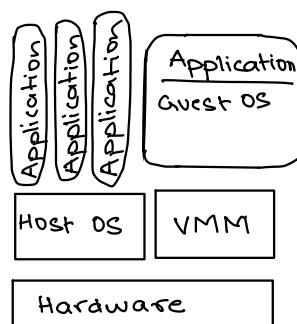
TRADITIONAL

- Thin layer of s/w that runs directly on the host machine's h/w
- Main adv. is performance



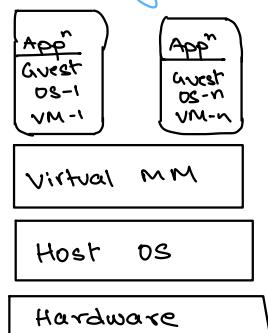
HYBRID

- Shares the h/w w/ the existing OS.



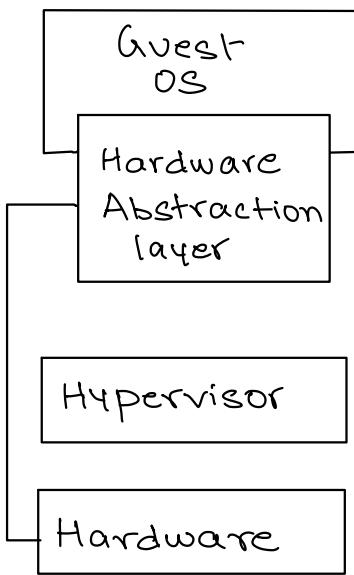
HOSTED

- The VM runs directly on top of the existing OS.



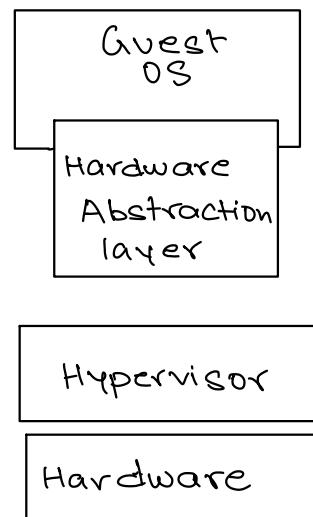
FULL VIRTUALIZATION

- Each VM runs on an exact copy of the actual hardware.
- Requires a fully virtualizable architecture; the hardware is fully exposed to the guest OS, that runs unchanged.



PARAVIRTUALIZATION

- VM runs on a modified copy of the actual hardware.
- It is done as some architectures are not easily virtualizable.
- Demands that the guest OS be modified to run under the VMM.



PROBLEMS FACED BY VIRTUALIZATION OF THE x86 ARCHITECTURE

① RING DEPRIVILEGING

VMM forces the guest SW, the OS, & the apps to run at a privilege level > 0^{not in (64 bit)}
solns → (0/1/3) mode
→ (0/3/3) mode

② RING RAISING

Problem when a guest OS is forced to run at a privilege level that it wasn't designed for.

③ ADDRESS SPACE COMPRESSION

VMM uses parts of the guest address space to store system DSs. These DS must be protected, but the guest SW must have access to them.

④ INTERRUPT VIRTUALIZATION

VMM generates a "virtual interrupt" in response to a physical interrupt & delivers it to the target guest OS. Guest OS's have the ability to mask interrupts
- complicates the VMM & increases overhead.

⑤ ACCESS TO HIDDEN STATE

Elements of the system state are hidden; there is no mechanism for saving & restoring hidden components after a context switch from 1 VM to another

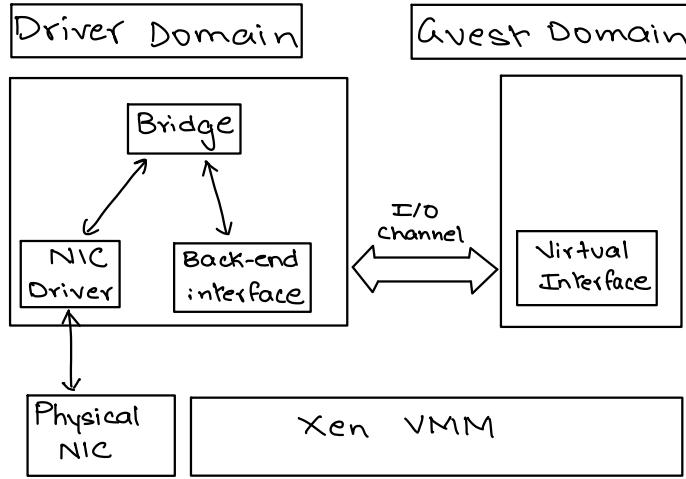
⑥ FREQUENT ACCESS TO PRIVILEGED RESOURCES INCREASES VMM OVERHEAD

The *TPR is frequently used by a guest OS. The VMM must protect access to this reg. & trap all attempts to access it. This can cause performance degradation.

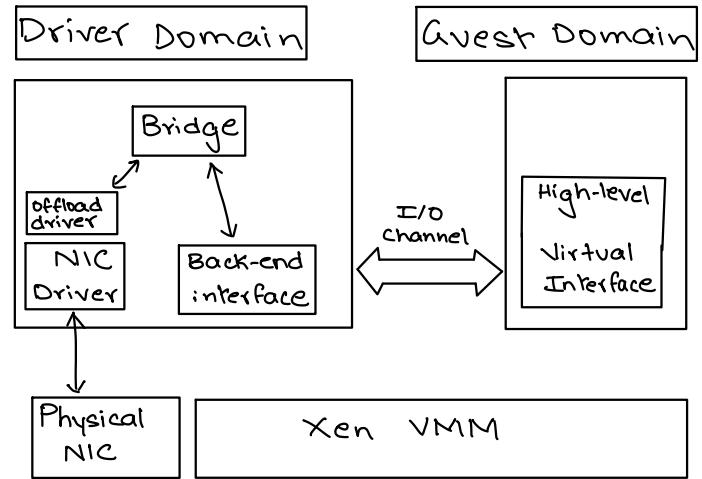
*TPR - Task Priority Register

XEN ARCHITECTURE

ORIGINAL



OPTIMIZED



XEN NETWORK OPTIMIZATIONS

① THE VIRTUAL INTERFACE

- The original "network provides" the guest domain w/ the abstraction of a simple low-level network interface supporting sending & receiving primitives.
- This design supports a wide range of physical devices attached to the driver domain but doesn't take advantage of the capabilities of some physical devices.
- These features are supported by the high-level virtual interface of the optimized system

② THE I/O CHANNEL

- Rather than copying a data buffer holding a packet, each packet is allocated in a new page & this physical page is remapped to the guest domain.
- This strategy contributes to a better than 4x increase in the send data rate.

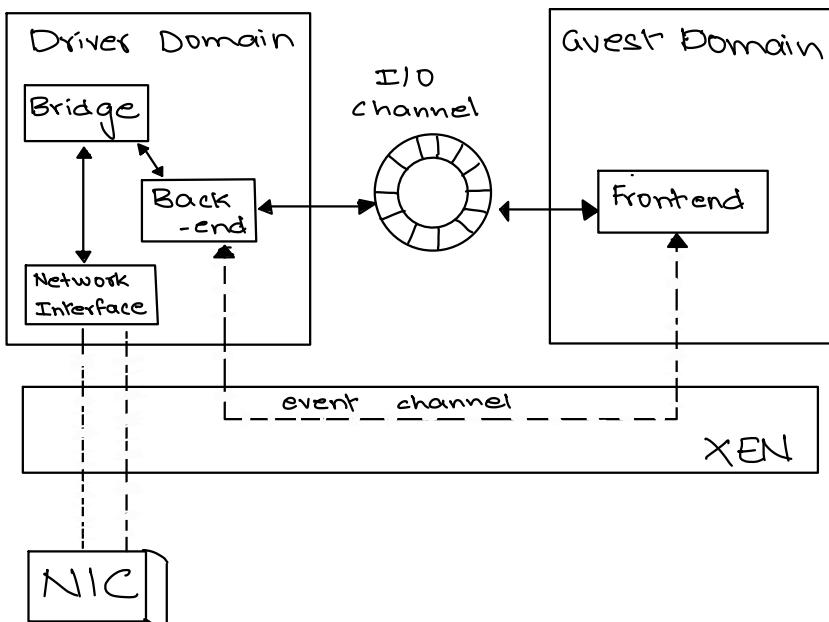
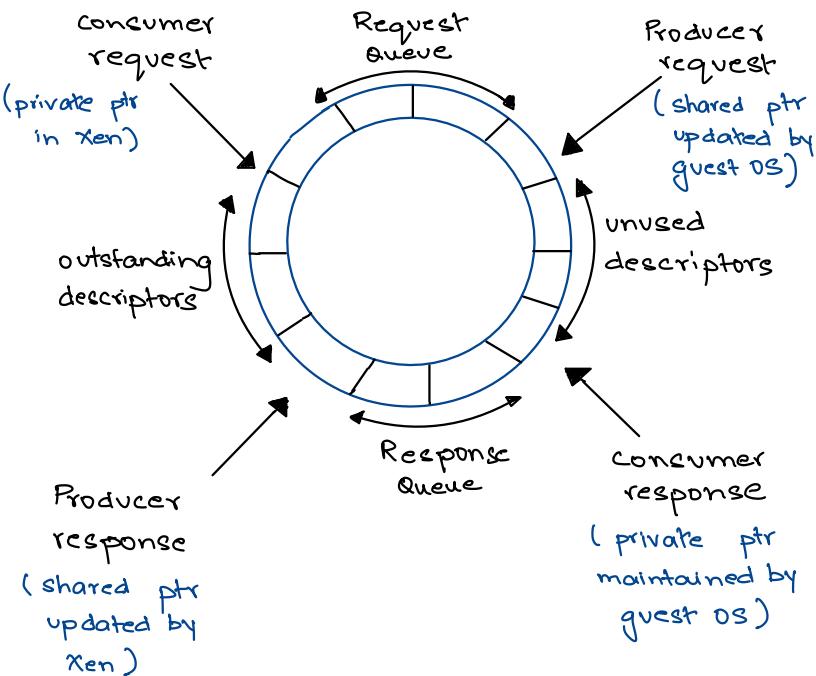
③ VIRTUAL MEMORY

- " in Xen 2.0 takes advantage of superpage & global page-mapping hardware features.
- A superpage increases the granularity of dynamic address translation - a superpage entry covers 1024 pages of physical memory, & address translation maps a set of contiguous pages to a set of contiguous physical pages.
- This reduces the no. of *TLB misses.
- The optimized version uses a special memory allocator to avoid the problem where the system is forced to use traditional page-mapping rather than superpage mapping

*Translation lookaside buffer: memory cache used to reduce the time taken to access a user memory location.

XEN ZERO-COPY SEMANTICS FOR DATA TRANSFER USING I/O RINGS

CIRCULAR RING OF BUFFERS



* IR → Instruction Register
 * RSE → Register Stack Engine

vBlades project

- The Itanium processor supports 4 privilege rings - PL0, PL1, PL2, PL3.
- Privileged instructions can only be executed by the kernel at PL0.
- Apps run at PL3 & can execute only non-privileged instructions.
- The VMM uses ring compression & runs itself at PL0 & PL1 while forcing a guest OS to run at PL2.
- A problem - privilege leaking - several nonprivileged instructions allow an appⁿ to determine the current privilege level (CPL).
- Itanium was selected because of its multiple functional units & multithreading support.

- CPU Virtualization:** when a guest OS attempts to execute a privileged instruction the VMM traps & emulates the instruction.

Complication - Itanium doesn't have an *IR & the VMM has to use state info to determine whether an instr. is privileged.

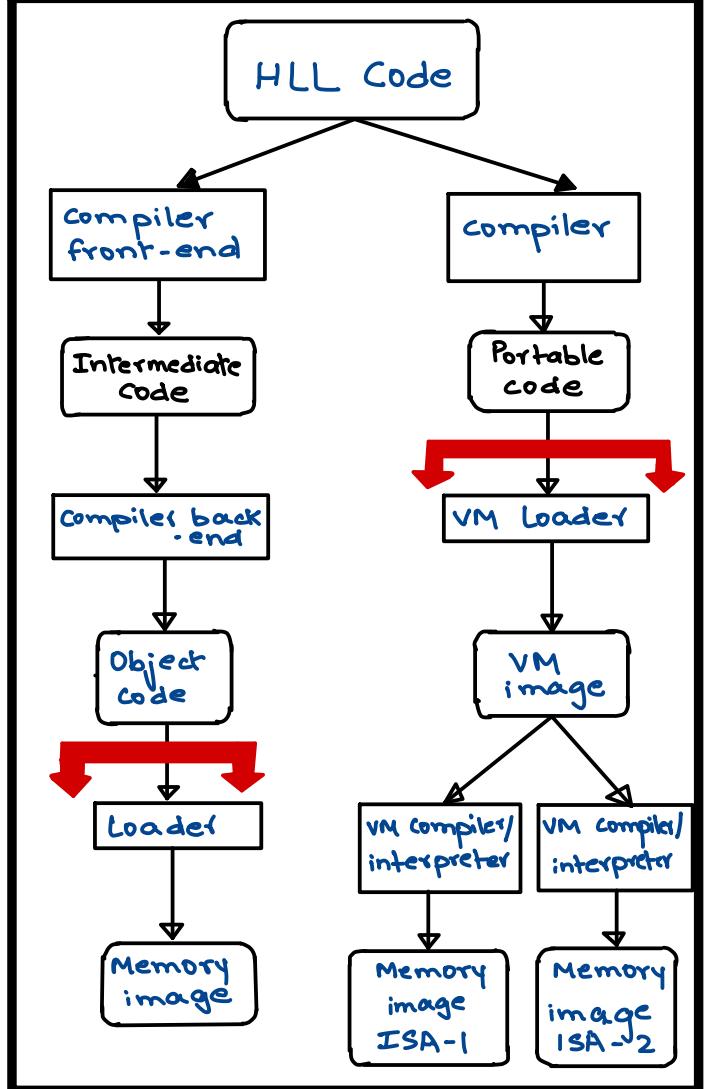
Complication - Caused by *RSE, which operates concurrently w/ the processor & may attempt to access memory & generate a page fault.

- A no. of privileged-sensitive instr. behave differently as a function of the privilege level. The VMM replaces each one of them w/ a privileged instr. during the dynamic transformation of the instruction stream.

contd.

contd.

- Memory virtualization is guided by the realization that a VMM should not be involved in most read & write operations to prevent a significant degradation of performance, but at the same time the VMM should exercise tight control & prevent a guest OS from acting maliciously.
- The vBlades VMM doesn't allow a guest OS to access memory directly. It inserts an additional layer of indirection called metaphysical addressing b/w virtual & real addressing.
- A guest OS is placed in metaphysical addressing mode. If the address is virtual, the VMM first checks if the guest OS is allowed to access that address, & if it is, it provides the regular address translation. If the address is physical, the VMM is not involved.



UNIT-4

CLOUD RESOURCE MANAGEMENT - POLICIES

① Admission Control

The goal of this policy is to prevent the system from accepting workloads in violation of high-level system policies — eg: a system may not accept an additional workload that would prevent it from completing work already in progress. Limiting the workload requires some knowledge about the global state of the system.

② Capacity Allocation

It means to allocate resources for individual instances; an instance is an activation of a service

③ Load Balancing, ④ Energy Optimization

- It refers to evenly distributing the load to a set of servers. In CC, a critical goal is minimizing the cost of providing the service, &, in particular, minimizing the energy consumption.
- This gives us a diff meaning of LB — instead of having the load equally distributed to all servers, we want to concentrate it & use the smallest no. of servers while switching the others to standby mode (server uses less energy).
- LB & energy optimization are correlated & affect the cost of providing the service.

⑤ QoS guarantees

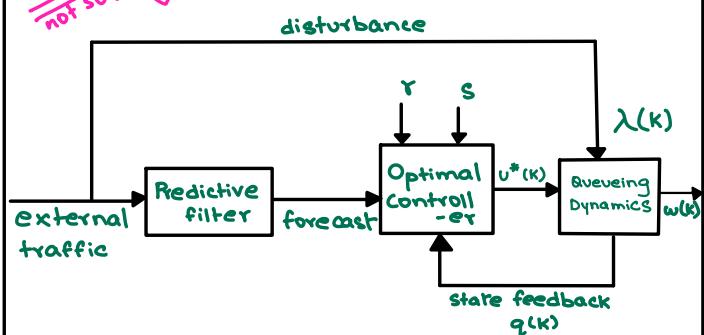
QoS is that aspect of resource management that is probably the most difficult to address &, at the same time, possibly the most critical to the future of CC.

(diagram)

- This is the case of a single processor serving a stream of input requests
- To compute the optimal inputs over a finite horizon, the controller uses feedback regarding the current state, as well as an estimation of the future disturbance due to the environment.

CONTROL THEORY PRINCIPLES — OPTIMAL RESOURCE ALLOCATION

not sure ↴



→ Optimal control generates a sequence of control inputs over a look-ahead horizon while estimating changes in operating conditions.

→ A convex cost function has arguments $x(k)$, the state at step k , & $u(k)$, the control vector; this cost function is minimized, subject to the constraints imposed by system dynamics.

→ The discrete-time optimal control problem is to determine the sequence of control variables $u(i), u(i+1), \dots, u(n-1)$ to minimize the expression:

$$J(i) = \underbrace{\phi(n, x(n))}_{\text{cost function of the final step}} + \sum_{k=i}^{n-1} \underbrace{L^k(x(k), u(k))}_{\text{Time-varying cost function at the intermediate step } k \text{ over } [i, n]}$$

→ The minimization is subject to :

$$x(k+1) = f^k(x(k), u(k))$$

\uparrow \uparrow \uparrow
 System state State at time $k+1$ Input at time k

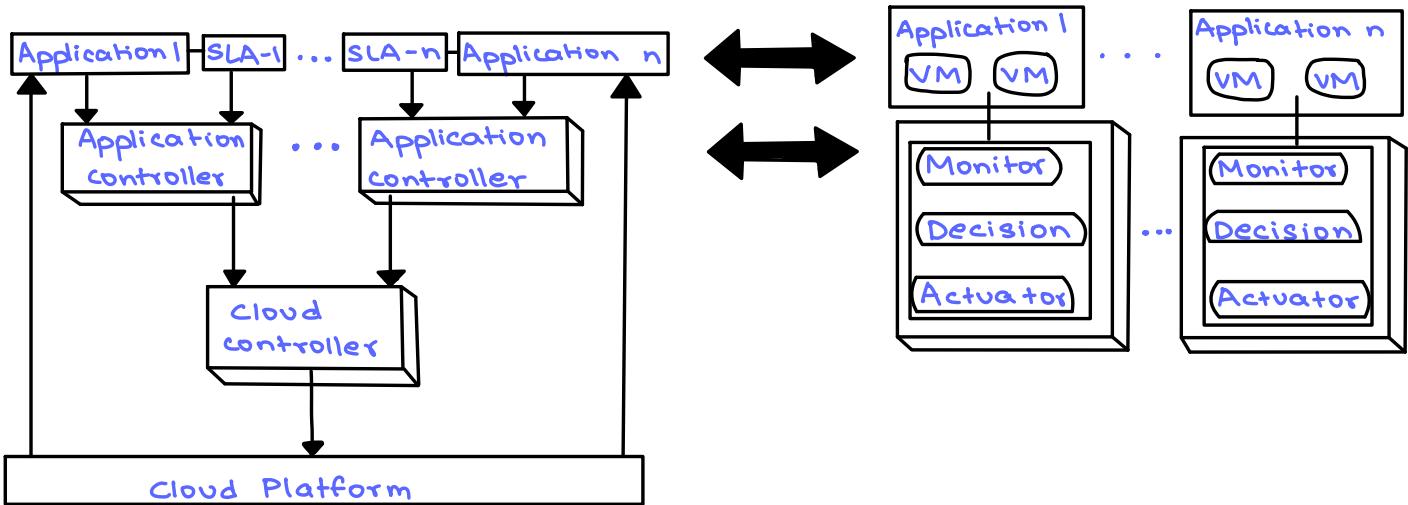
→ A technique to solve this problem
— Lagrange multiplier (λ) method to maximize a function $g(x, y)$ subject to constraint $h(x, y) = K$

$$\Lambda(x, y, \lambda) = g(x, y) + \lambda \times [h(x, y) - K]$$

→ A necessary condition for optimality is that (x, y, λ) is a stationary point for $\Lambda(x, y, \lambda)$, or,

$$\nabla_{x, y, \lambda} \Lambda(x, y, \lambda) = 0$$

TWO-LEVEL CONTROL ARCHITECTURE



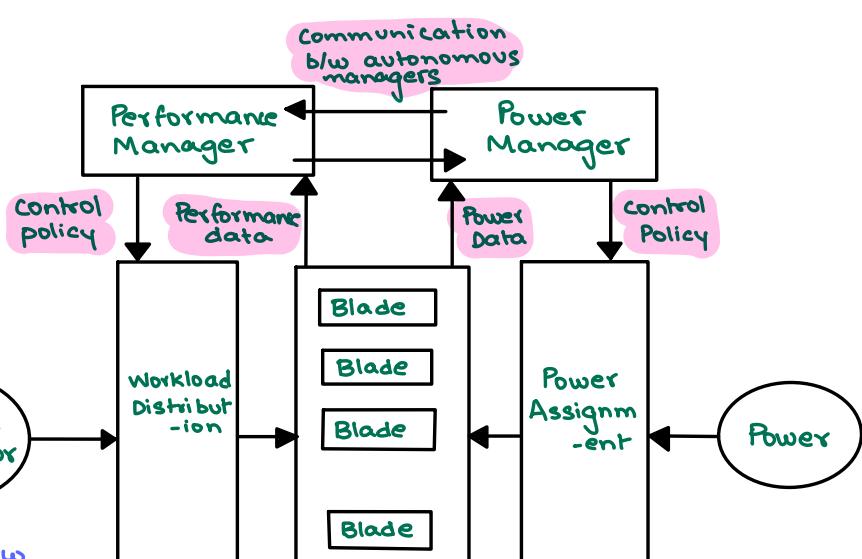
- The automatic resource management is based on 2 levels of controllers - one for the service provider & one for the application.
- The main components of a control system are the inputs, the control system components, & the outputs.
- The inputs in such models are the offered workload & the policies for admission control, the capacity allocation, the load balancing, the energy optimization, & the QoS guarantees in the cloud.
- The system components are sensors used to estimate relevant measures of performance & controllers that implement various policies.
- The output is the resource allocations to the individual applications.
- The controllers use the feedback provided by sensors to stabilize the system; stability is related to the change of the output.
- If the change is too large, the system may become unstable.

COORDINATION OF AUTONOMIC PERFORMANCE MANAGERS

*The approach to coordinating power & performance management is based on several ideas:

- Use a joint utility function for power & performance. This function, $U_{pp}(R, P)$ is a function of the response time R , & the power P .

$$U_{pp}(R, P) = \frac{U(R)}{P}$$



- Identify a minimal set of parameters to be exchanged b/w the 2 managers.
- Set up a power cap for individual systems based on the utility-optimized power management policy
- Use a std performance manager modified only to accept input from the power manager regarding the frequency determined according to

contd.

- the power management policy.
- Use standard software systems.

utility function $U'(P_k, n_c) = U_{pp}(R(P_k, n_c), P(P_k, n_c))$

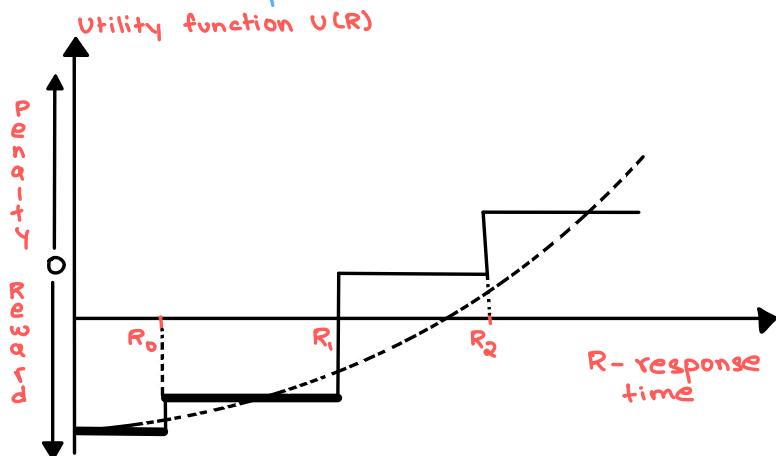
↑
power cap ↑
No. of clients

UTILITY FUNCTION - PERFORMANCE METRIC : RESPONSE TIME (R)

- A utility function relates the "benefits" of an activity or service w/ the "cost" to provide the service
- A service-level-agreement (SLA) often specifies the rewards & penalties associated w/ specific performance metrics.

(R₀, R₁, R₂)

When the performance metric is R:



- The largest reward can be obtained when $R \leq R_0$; a slightly lower reward corresponds to $R_0 < R \leq R_1$.
- when $R_1 < R < R_2$, instead of gaining a reward, the provider of the service pays a small penalty; the penalty increases when $R > R_2$.
- A utility function, U(R), which captures this behavior, is a sequence of step functions.
- The utility function is sometimes approximated by a quadratic curve.

PRICING & ALLOCATION ALGORITHM

This algo partitions the set of users into 2 disjoint sets — winners (W) & losers (L).

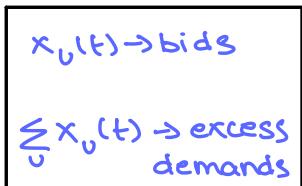
The algo should:

- ★ Be computationally tractable.
- ★ Scale well
- ★ Be objective. Partitioning should only be based on the price Π_u of a user's bid.
- ★ Be fair. Make sure that the prices are uniform. All winners w/in a given resource pool pay the same price.
- ★ Indicate clearly at the end of the auction the unit prices for each resource pool.
- ★ Indicate clearly to all participants the relationship b/w the supply & the demand in the system

ACSA COMBINATORIAL AUCTION ALGORITHM

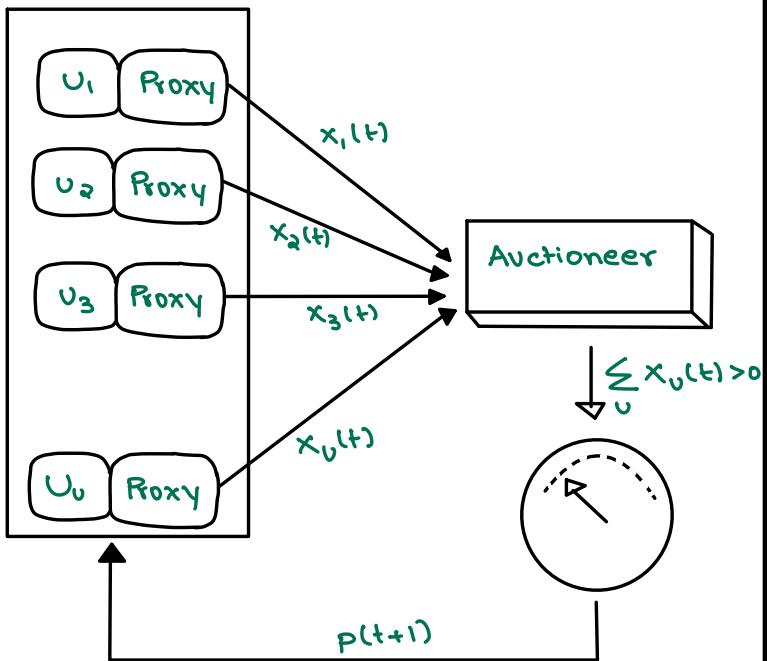
- Here, the participants at the auction specify the resource & the quantities of that resource offered or desired at the price listed for that time slot.
- Then the excess vector is calculated. If all its components are -ve, the auction stops. (Demand doesn't exceed the offer).
- If Demand > Offer, the auctioneer increases the price for items & solicits bids at the new price.

Schematics:



To allow for a single round, auction users are represented by proxies that place the bids $x_v(t)$.

The auctioneer determines whether there is an excess demand, & in that case, raises the price of those resources.



FAIR QUEUEING

→ It introduces a bit-by-bit round-robin (BR) strategy; a single bit from each queue is transmitted & the queues are visited in a round-robin fashion.

$R(t)$ → No. of rounds of the BR algo up to time t

$N_{\text{active}}(t)$ → No. of active flows through the switch

t_i^a → Time when packet i of flow a , of size P_i^a arrives

S_i^a → Value of $R(t)$ when the 1st bit of packet i of flow a is transmitted

F_i^a → Value of $R(t)$ when the last bit of packet i of flow a is transmitted

then,

$$F_i^a = S_i^a + P_i^a$$

$$S_i^a = \max[F_{i-1}^a, R(t_i^a)]$$

→ The quantities $R(t)$, $N_{\text{active}}(t)$, S_i^a & F_i^a depend only on the arrival time of the packets, t_i^a , & not on their transmission time, provided that a flow a is active as long as:

$$R(t) \leq F_i^a, \quad \text{when } i = \max(j \mid t_j^a \leq t)$$

→ The next packet to be transmitted is the one w/ the smallest F_i^a .

START-TIME FAIR QUEUEING

* The basic idea is to organize the consumers of the CPU bandwidth in a tree-structure; the root node is the processor & the leaves are the threads of each application.

* A scheduler acts at each level of the hierarchy. The fraction of the processor bandwidth, B , allocated to the intermediate node i is:

$$\frac{B_i}{B} = \frac{w_i}{\sum_{j=1}^n w_j}$$

weight of the
n children of the
node i

Rules an SFQ scheduler follows:

R1: The threads are serviced in the order of their virtual start-up time

R2: The virtual start-up time of the i -th activation of thread x is:

$$S_x^i(t) = \max[v(\tau^i), F_x^{(i-1)}(t)] \quad \& \quad S_x^0 = 0$$

The condition for thread i to be started is that thread $(i-1)$ has

finished & that the scheduler is active.

R3: The virtual finish time of the i -th activation of thread x is

$$F_x^i(t) = S_x^i(t) + \frac{q}{w_x} \quad \begin{matrix} \text{Time} \\ \text{quantum of} \\ \text{the thread} \end{matrix}$$

A thread is stopped when its time quantum has expired.

R4: The virtual time of all threads is initially zero,

$$v_x^0 = 0$$

The virtual time $v(t)$ at real time t is computed:

$$v(t) = \begin{cases} \text{Virtual start time of the} \\ \text{thread in service at} \\ \text{time } t & \text{if CPU} \\ & \text{is busy} \\ \text{Max finish virtual} \\ \text{time of any thread} & \text{if CPU} \\ & \text{is idle} \end{cases}$$

SUM - SFQ

Q. There are 2 threads w_1 weights $w_a = 1$ & $w_b = 4$. Time quantum $q = 12$. (Consider scheduling decisions at $t=0$ & $t=3$)

Initially, $S_a^0 = 0$, $S_b^0 = 0$
 $v_a(0) = 0$, $v_b(0) = 0$

① $t=0$

$$S_a^0 = S_b^0 \Rightarrow \text{tie}$$

Thread b is arbitrarily chosen to run first.

The virtual finish time of thread b is:

$$F_b^0 = S_b^0 + \frac{q}{w_b}$$

$$= 0 + \frac{12}{4} = 3$$

(pg 187-189 in the tb
for more values
of t)

② $t=3$

Both threads are runnable & thread b was in service; thus, $v(3) = S_b^0 = 0$; then

$$S_b^1 = \max[v(3), F_b^0]$$

$$= \max(0, 3)$$

$$= 3$$

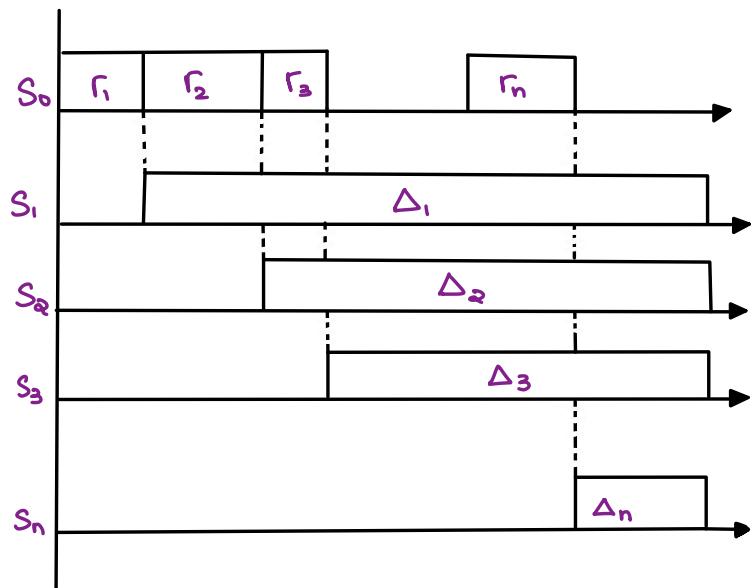
But $S_a^0 < S_b^1$, thus thread a is selected to run. Its virtual finish time:

$$F_a^0 = S_a^0 + \frac{q}{w_a}$$

$$= 0 + \frac{12}{1} = 12$$

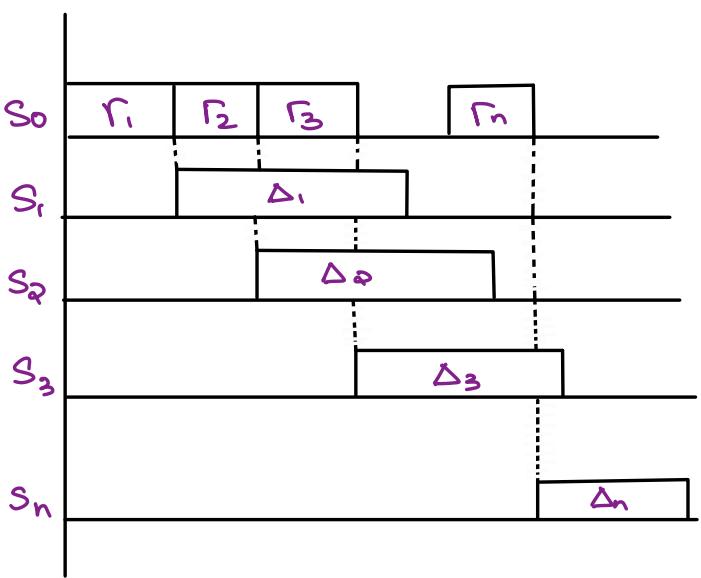
OPTIMAL PARTITIONING RULE

- The workload is partitioned to ensure the earliest possible completion time.
- The algorithm requires worker nodes to complete to complete execution at the same time.
- The head node, S_0 , distributes sequentially the data to individual worker nodes.



EQUAL PARTITIONING RULE

- EPR assigns an equal workload to individual worker nodes.
- The head node, S_0 , distributes sequentially the data to individual worker nodes.



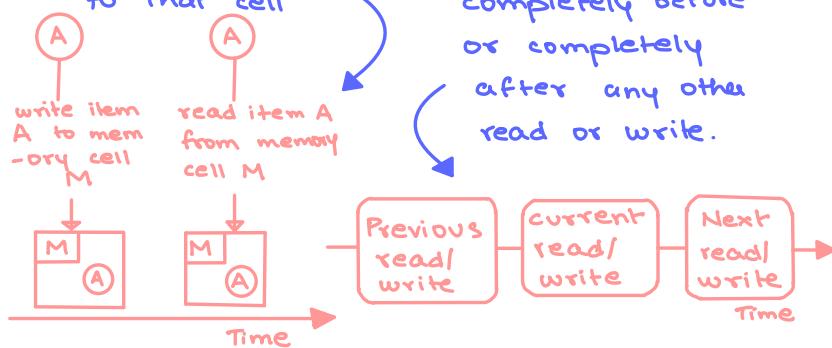
UNIT-5

TWO ABSTRACT MODELS OF STORAGE

① CELL STORAGE

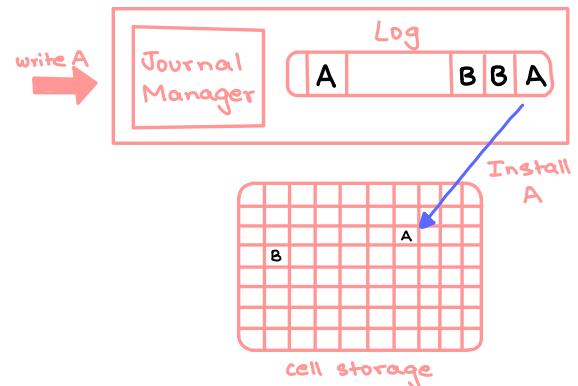
- Assumes that the storage consists of cells of the same size & that each object fits exactly in one cell.
- This model reflects the physical organization of several storage media; a secondary storage device is organized in sectors or blocks read & written as a unit.
- read/write coherence** & **before-or-after atomicity** are 2 highly desirable properties of any storage model & in particular of cell storage.

The result of a read of memory cell M should be the same as the most recent write to that cell



② JOURNAL STORAGE

- It is a fairly elaborate organization for storing composite objects such as records consisting of multiple fields.
- Consists of a manager & cell storage, where the entire history of a variable is maintained, rather than just the current value.
- A log contains a history of all variables in cell storage.
- The user doesn't have direct access to the cell storage; instead the user can request the journal manager to,
 - (i) start a new action
 - (ii) read the value of a cell,
 - (iii) write "
 - (iv) commit an action , (v) abort an "
- The journal manager translates user requests to commands sent to the cell storage :
 - (i) read a cell, (ii) write " "
 - (iii) allocate " ", (iv) deallocate " "

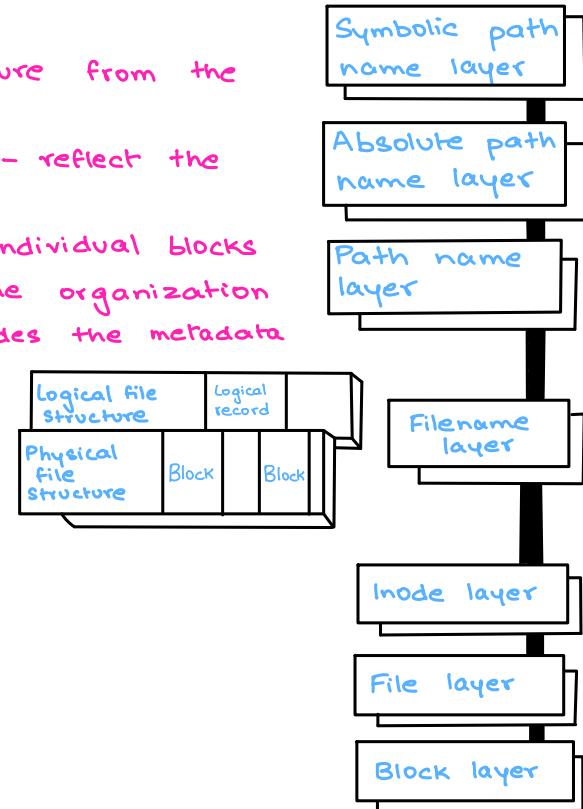


CHARACTERISTICS OF UFS (UNIX FILE SYSTEM)

- The layered design provides the necessary **flexibility** for the file system; **layering** allows separation of concerns & minimization of the interaction among the modules. The addition of vnode layer allowed the UFS to treat local & remote file access uniformly.
- The **hierarchical design** supports **scalability** of the file system; it allows grouping of files into directories & supports multiple levels of directories.
- The **metadata** supports a **systematic** rather than an ad hoc design philosophy. inodes contain info about individual files & directories. Metadata includes the file owner, access rights, creation time or time of last modif., file size & storage device cells.

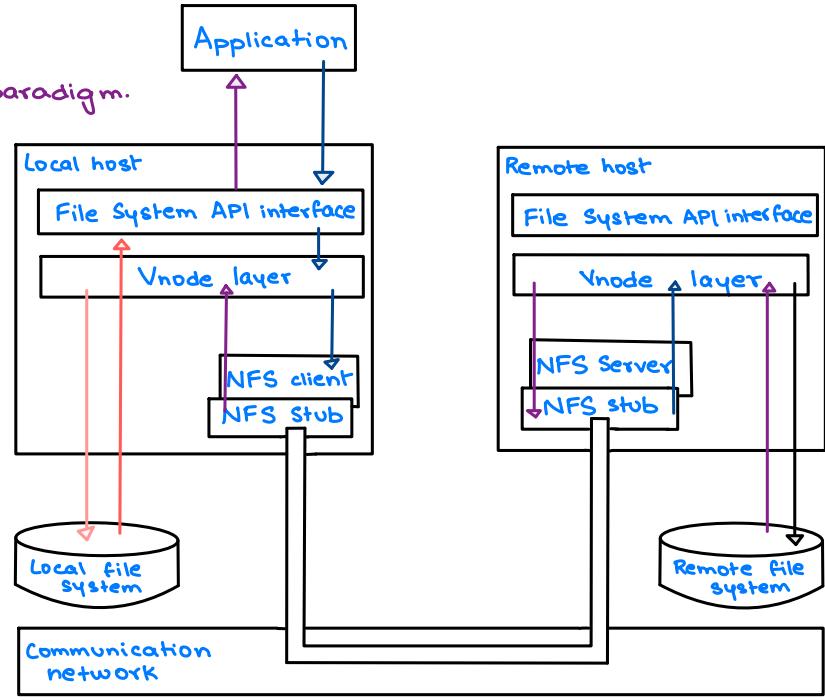
LAYERED DESIGN OF UFS

- This design separates the physical file structure from the logical one.
- The lower 3 layers - the block, file & inode layer - reflect the physical organization.
- The block layer allows the system to locate individual blocks on the physical device; the file layer reflects the organization of blocks into files; & the inode layer provides the metadata for the objects.
- The upper 3 layers reflect the logical organization.
- The filename layer mediates b/w the machine-oriented & the user-oriented views of the file system.

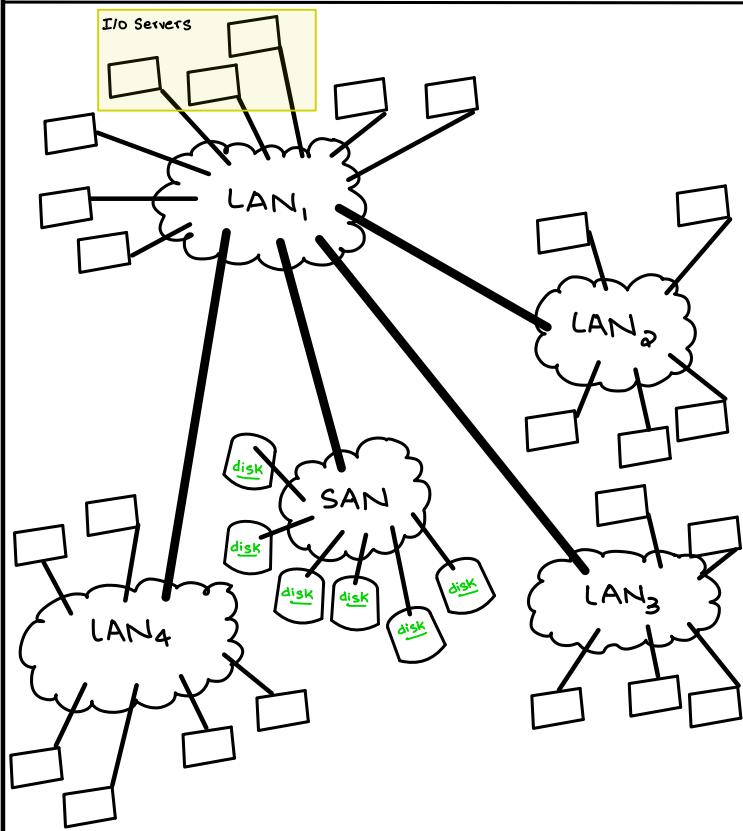


NETWORK FILE SYSTEM

- NFS is based on a client-server paradigm.
- The client runs on the local host while the server is at the site of the remote file system, & they interact by means of remote procedure calls (RPCs).
- The API interface of the local file system distinguishes file operations on a local file from the ones on a remote file, & in the latter case, invokes the RPC client.
- NFS uses a Vnode layer to distinguish b/w operations on local & remote files.



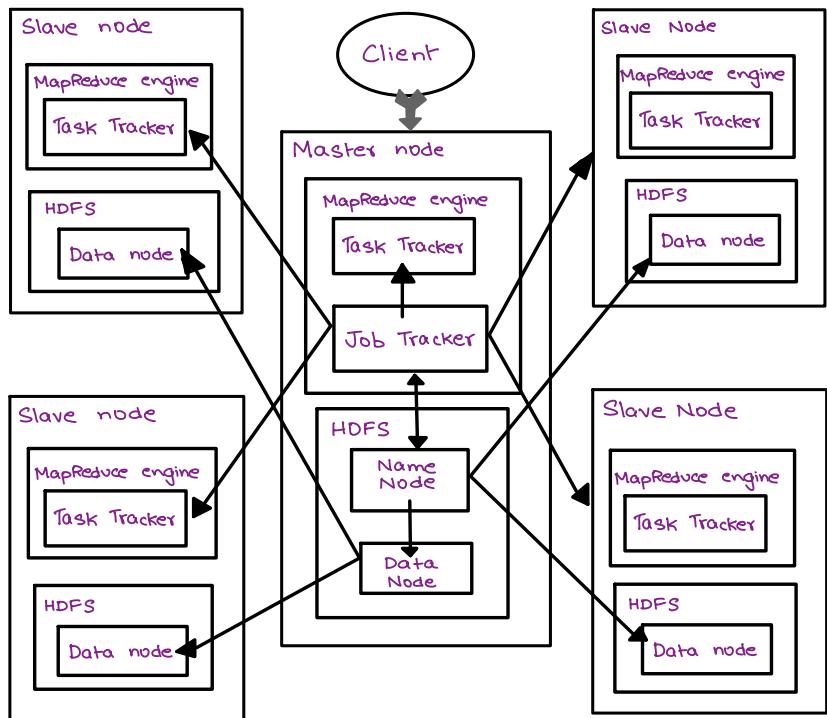
GENERAL PARALLEL FILE SYSTEM



- * Parallel file systems allow multiple clients to read & write concurrently from the same file.
- * GPFS is a parallel file system that emulates closely the behavior of a general-purpose POSIX system running on a single system.
- * It was designed for optimal performance of large clusters.
- * Reliability is a major concern in a system w/ many physical components.
- * To recover from system failure, GPFS records all metadata updates in a write-ahead log file.
- * The log files are maintained by each I/O node for each file system it mounts.
- * GPFS uses disk maps to manage disk space

HADOOP - HDFS

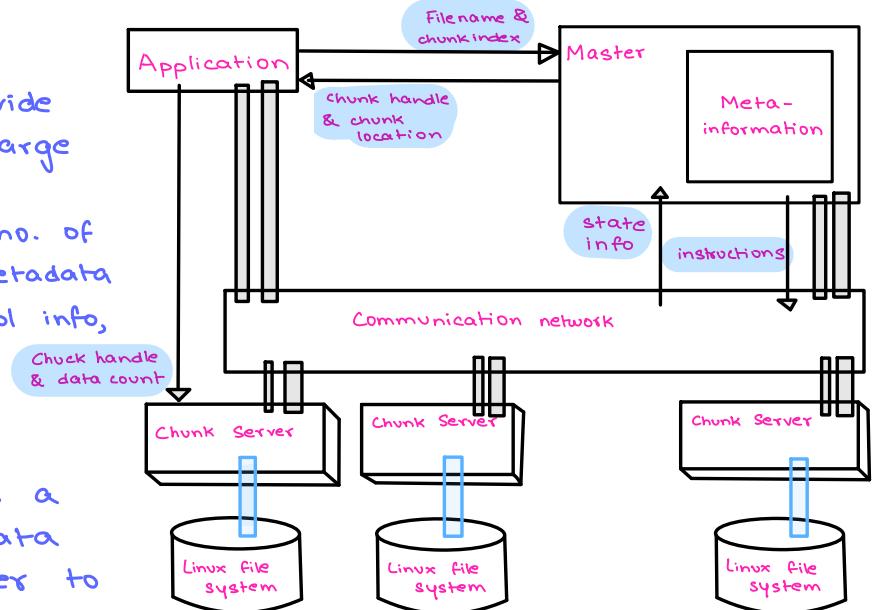
- Hadoop supports distributed app's handling extremely large volumes of data.
- A Hadoop system has 2 components — a MapReduce engine & a database. The database could be HDFS, Amazon S3, or cloudStore.
- The Hadoop engine on the master of a multinode cluster consists of a job tracker & a task tracker, whereas the engine on a slave has only a task tracker.
- The job tracker receives a MapReduce job from a client & dispatches the work to the task trackers running on the nodes of a cluster.
- The task tracker supervises the execution of the work allocated to the node.
- HDFS replicates data on multiple nodes — the default is 3 replicas.
- The name node running on the master manages the data distribution & data replication & communicates w/ data nodes running on all cluster nodes.



GOOGLE FILE SYSTEM (GFS)

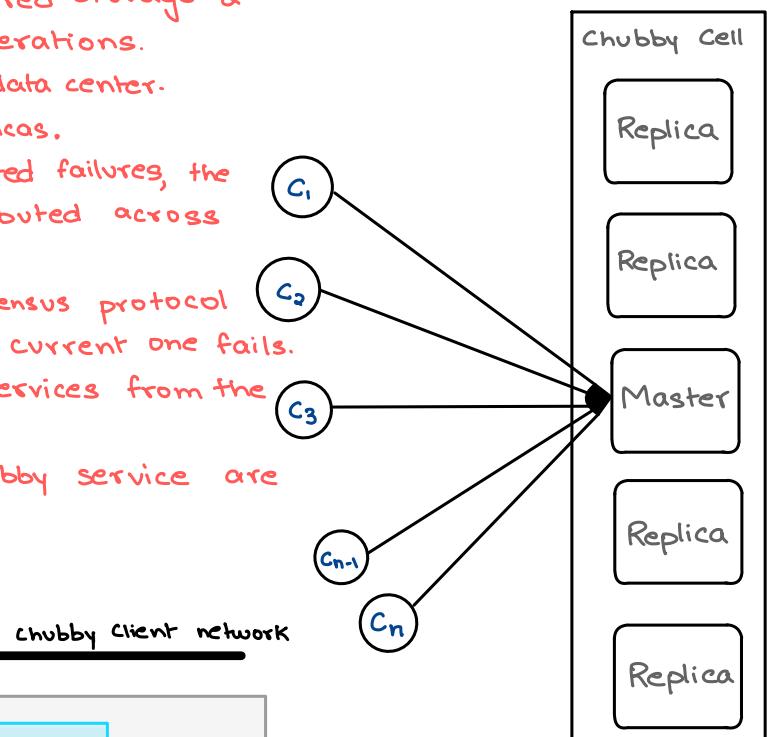
- * It uses thousands of storage systems built from inexpensive commodity components to provide petabytes of storage to a large user community.
- * A master controls a large no. of chunk servers; it maintains metadata such as filenames, access control info, state of individual chunk servers & the location of all the replicas for every chunk.
- * The operation log maintains a historical record of metadata changes, enabling the master to recover in case of a failure.
- * Each chunk server is a commodity Linux system; it receives instructions from the master & responds with status info.
- * To access a file, an app sends to the master the filename & the chunk index, the offset in the file for the read/write operation; the master responds w/ the chunk handle & the location of the chunk.

Architecture of a GFS cluster



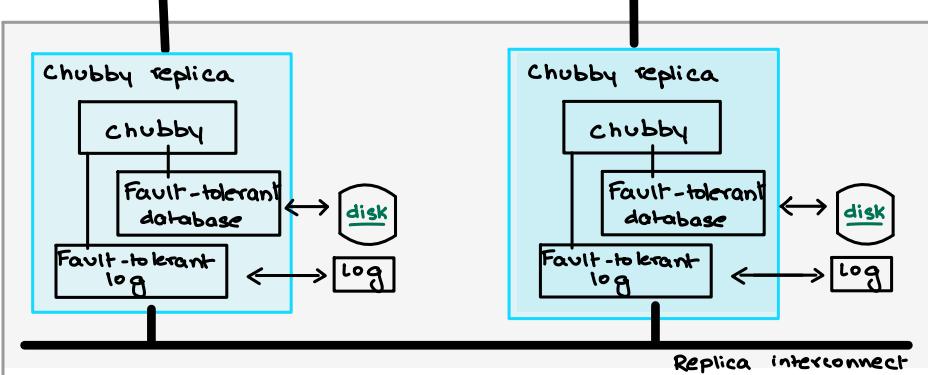
LOCKS & CHUBBY

- Locks enable controlled access to shared storage & ensure atomicity of read & write operations.
- A chubby cell typically serves 1 data center.
- The cell server includes several replicas.
- To reduce the probability of correlated failures, the servers hosting replicas are distributed across the campus of a data center.
- The replicas use a distributed consensus protocol to elect a new master when the current one fails.
- Clients use RPCs to request services from the master.
- The files & directories of the chubby service are organized in a tree structure.



Chubby replica architecture

chubby client network

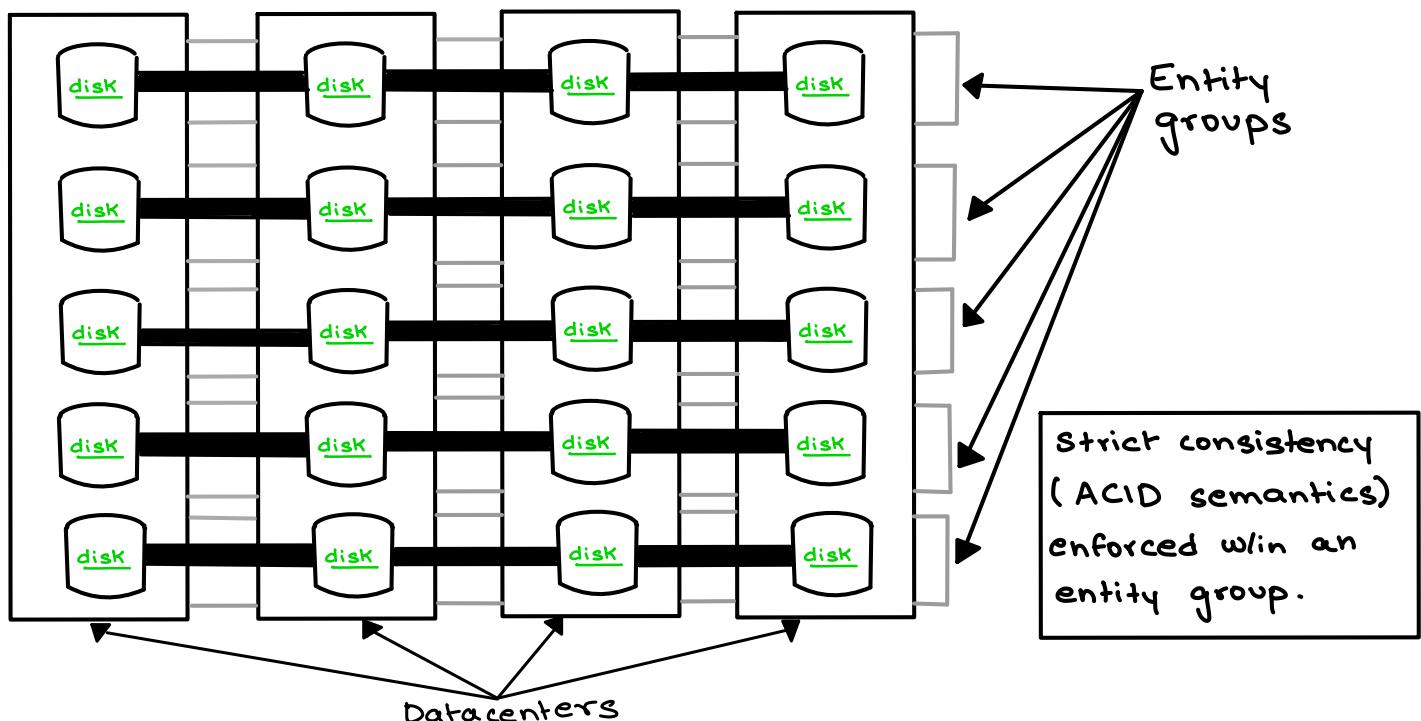


PAXOS ALGO

- It is used to reach consensus on sets of values. To ensure that the Paxos algo succeeds in spite of the occasional failure of a replica, the following 3 phases of the algo are executed repeatedly.
- ① Elect a replica to be the master/coordinator.
- ↳ When a master fails, several replicas may decide to assume the role of a master.
 - ↳ To ensure that the result of the election is unique, each replica generates a sequence no. larger than any sequence no. it has seen & broadcasts it in a purpose message.
 - ↳ The replicas that haven't seen a higher sequence no. broadcast a promise reply & declare that they will reject proposals from other candidate masters.
 - ↳ If the no. of respondents represents a majority of replicas, the one that sent the purpose msg is elected master.
- ② The master broadcasts to all replicas an accept message, including the value it has selected, & waits for replies — acknowledge/reject.
- ③ Consensus is reached when the majority of the replicas send an acknowledge msg; then the master broadcasts the commit msg.

MEGASTORE

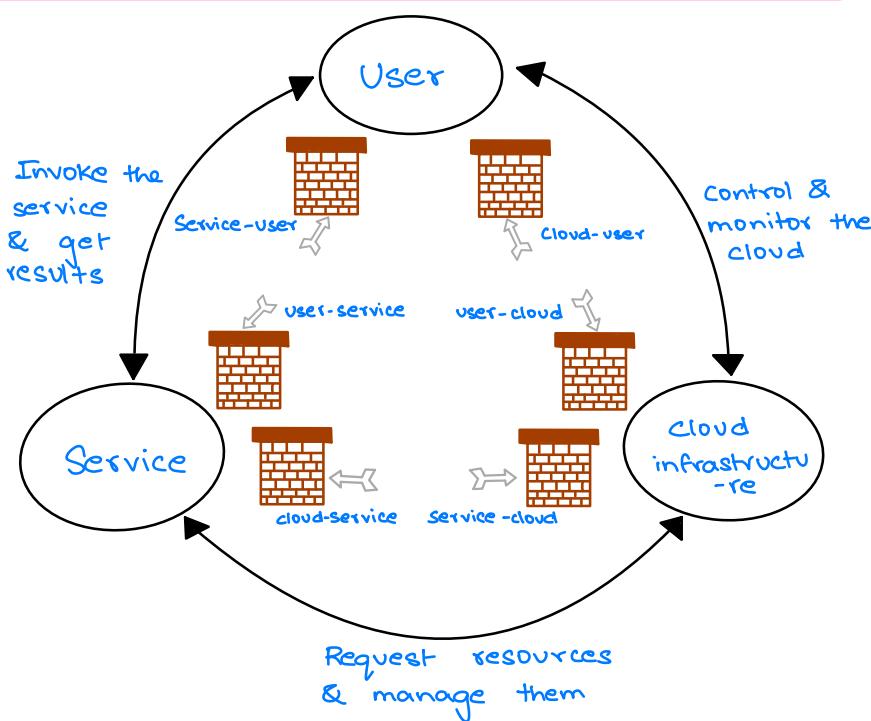
- It is scalable storage for online services.
- The system, distributed over several data centers, has a very large capacity & it is highly available.
- The basic design philosophy of the system is to partition the data into entity groups & replicate each partition independently in data centers located in different geographic areas.
- Megastore supports only those traditional DB features that allow the system to scale well & that don't drastically affect the response time.
- It uses Paxos consensus algo to replicate primary user data, metadata, & system configuration info across data centers & for locking.
- The entity groups are app-specific & store together logically related data.
- The data model is declared in a schema consisting of a set of tables composed of entries, each entry being a collection of named & typed properties.
- A Megastore table can be a root or a child table.



CLOUD SECURITY RISKS {

- Classes of risk — traditional threats, threats related to system availability, & threats related to 3rd-party data control.
- Traditional threats begin at the user site. The user must protect the infrastructure used to connect to the cloud.
- The next threat is related to the authentication & authorization process.
- Moving to the cloud, there are many types of traditional attacks — DDoS, phishing, SQL injection.
- Availability of cloud services is another concern. System failures, power outages & other catastrophic events could shut down cloud services for long periods of time.
- 3rd-party control generates concerns caused by the lack of transparency & limited user control.

SURFACE OF ATTACKS IN A CC ENVIRONMENT



- The 3 actors involved in the model considered are — the user, the service, & the cloud infrastructure.
- There are 6 types of attacks possible.
- The User can be attacked from 2 directions: from the service & from the cloud.
 - Spoofting that originates from the cloud infrastructure
 - SSL certificate spoofing, attacks on browser caches, phishing
- The service can be attacked from the user & from the cloud.
 - Buffer overflow, SQL injection, privilege escalation
 - Limiting access to resources, privilege-related attacks, data distortion, injecting additional operations
- The cloud infrastructure can be attacked by a user who targets the cloud control system.
- The " " may also be targeted by a service requesting an excessive amount of resources & causing the exhaustion of resources.

SECURITY, PRIVACY & TRUST

① SECURITY

- Security is the top concern for cloud users, who are accustomed to having full control of all systems on which sensitive info is stored & processed.
- Major user concerns are unauthorized access to confidential info & data theft.
- Data is more vulnerable in storage than while it is being processed.
- Threats during processing can originate from flaws in the VMM, rogue VMs, or a VMBR.
- The next concerns regard user control over the life cycle of data. It is virtually impossible for a user to determine whether data that should have been deleted is actually deleted.
- Lack of standardization is another concern. Today there are no interoperability standards.

② PRIVACY

- The term privacy refers to the right of an individual, a group of individuals, or an organization to keep info of a personal or proprietary nature from being disclosed to others.
- Our discussion targets primarily public clouds where data, often in an encrypted form, resides on servers owned by a CSP.
- The owner of the data cannot rely exclusively on the CSP to guarantee the privacy of data.

- The main aspects of privacy are - the lack of user control, potential unauthorized secondary use, data proliferation, & dynamic provisioning.
- A CSP may obtain revenues from unauthorized secondary usage of info, eg - for targeted marketing.

③ TRUST

- It means "assured reliance on the character, ability, strength, or truth of someone or something".
- 2 conditions must exist for trust to develop:
 - ↳ Risk - the perceived probability of loss.
 - ↳ Interdependence - the idea that the interests of 1 entity cannot be achieved without reliance on other entities.
- A trust relationship goes through 3 phases:
 - (1) a building phase - trust is formed.
 - (2) a stability " - when trust exists.
 - (3) a dissolution " - " declines
- The Internet offers individuals the ability to obscure or conceal their identities.
- The resulting anonymity reduces the cues normally used in judgements of trust.
- Policies & reputation are 2 ways of determining trust.

THREATS IDENTIFIED BY THE NIST PROJECT

VMM-Based threats :

① Starvation of resources & denial of service for some VMs.

Probable causes - badly configured resource limits, rogue VM w/ the capability to bypass resource limits set in the VMM.

② VM side-channel attacks by a rogue VM under the same VMM.

Probable causes - lack of proper isolation of inter-VM traffic due to misconfiguration, limitation of packet inspection devices, presence

of VM instances built from insecure VM images.

③ Buffer overflow attacks.

VM-Based threats :

① Deployment of rogue or insecure VM.

Probable cause - improper config of access controls on VM admin tasks.

② Presence of insecure & tampered VM images in the VM image repository

Probable causes - lack of access control to the VM image repository, lack of mechanisms to verify the integrity of the images.