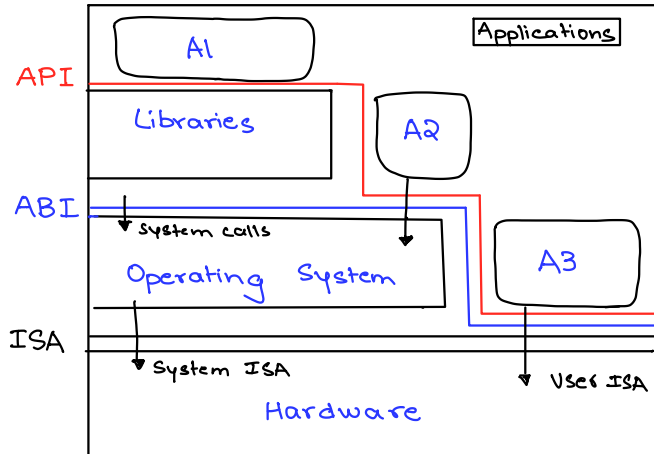




UNIT 3

LAYERING

- managing system complexity - layering
- minimizes interactions
- 2 modes → user & privileged




- 1st interface - ISA - boundary b/w s/ware & h/ware - processor's set of instructions
- ABI (Appⁿ Binary Int.) - H/ware access to the appⁿs & library modules. Doesn't include privileged instructions.
- API - defines set of instr the h/ware was designed to execute - gives appⁿ access to the ISA.

CONDITIONS FOR EFFICIENT VIRTUALIZATION

- ① A program running under the VMM must exhibit behaviour essentially similar to that demonstrated when the appⁿ runs directly on an equivalent machine.
- ② The VMM should be in complete control of the virtualized resources.
- ③ A statistically significant fraction of the machine instructions should be executed without the intervention of the VMM.


TRADITIONAL

- Thin layer of s/ware that runs directly on the host machine's h/ware
- Main adv. is performance




HYBRID

- Shares the h/ware w/ the existing OS.




HOSTED

- The VM runs directly on top of the existing OS.




FULL VIRTUALIZATION

- Each VM runs on an exact copy of the actual hardware.
- Requires a fully virtualizable architecture; the hardware is fully exposed to the guest OS, that runs unchanged.



PARAVIRTUALIZATION

- VM runs on a modified copy of the actual hardware.
- It is done as some architectures are not easily virtualizable.
- Demands that the guest OS be modified to run under the VMM.



PROBLEMS FACED BY VIRTUALIZATION OF THE x86 ARCHITECTURE

① RING DEPRIVILEGING

VMM forces the guest software, the OS, & the apps to run at a privilege level > 0 ^{not in (64 bit)} _{solns → (0/1/3) mode}
→ (0/3/3) mode

② RING ABUSING

Problem when a guest OS is forced to run at a privilege level that it wasn't designed for.

③ ADDRESS SPACE COMPRESSION

VMM uses parts of the guest address space to store system DSs. These DS must be protected, but the guest software must have access to them.

④ INTERRUPT VIRTUALIZATION

VMM generates a "virtual interrupt" in response to a physical interrupt & delivers it to the target guest OS. Guest OSs have the ability to mask interrupts - complicates the VMM & increases overhead.

⑤ ACCESS TO HIDDEN STATE

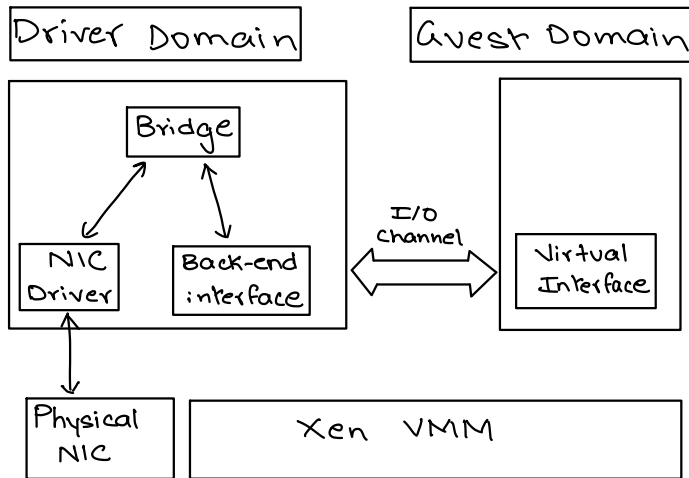
Elements of the system state are hidden; there is no mechanism for saving & restoring hidden components after a context switch from 1 VM to another.

⑥ FREQUENT ACCESS TO PRIVILEGED RESOURCES INCREASES VMM OVERHEAD

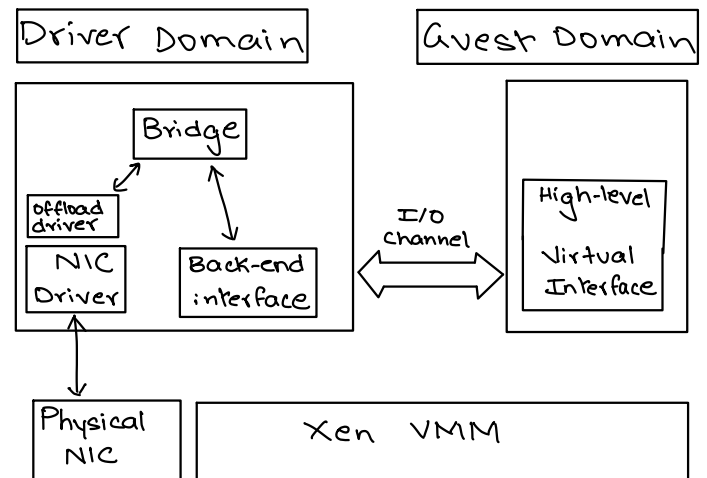
The ^{*}TPR is frequently used by a guest OS. The VMM must protect access to this reg. & trap all attempts to access it. This can cause performance degradation.

XEN ARCHITECTURE

ORIGINAL



OPTIMIZED



XEN NETWORK OPTIMIZATIONS

① THE VIRTUAL INTERFACE

- The original " " network provides the guest domain w/ the abstraction of a simple low-level network interface supporting sending & receiving primitives.
- This design supports a wide range of physical devices attached to the driver domain but doesn't take advantage of the capabilities of some physical devices.
- These features are supported by the high-level virtual interface of the optimized system

② THE I/O CHANNEL

- Rather than copying a data buffer holding a packet, each packet is allocated in a new page & this physical page is remapped to the guest domain.
- This strategy contributes to a better than 4x increase in the send data rate.


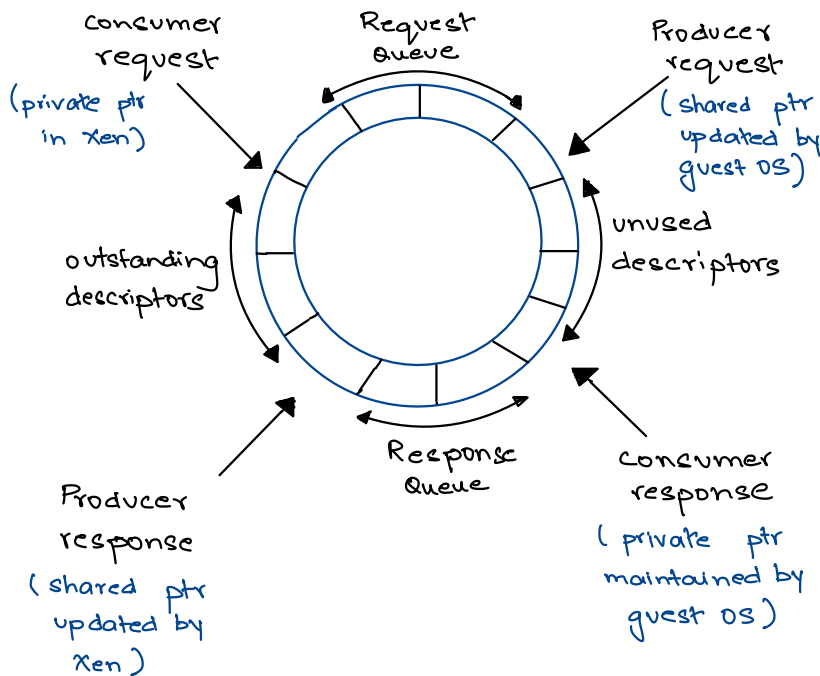
③ VIRTUAL MEMORY

- " " in Xen 2.0 takes advantage of superpage & global page-mapping hardware features.
- A superpage increases the granularity of dynamic address translation - a superpage entry covers 1024 pages of physical memory, & address translation maps a set of contiguous pages to a set of contiguous physical pages.
- This reduces the no. of *TLB misses.
- The optimized version uses a special memory allocator to avoid the problem where the system is forced to use traditional page-mapping rather than superpage mapping

*Translation lookaside buffer:
memory cache used to
reduce the time taken to
access a user memory
location.

XEN ZERO-COPY SEMANTICS FOR DATA TRANSFER USING I/O RINGS

CIRCULAR RING OF BUFFERS



* IR → Instruction Register
* RSE → Register Stack Engine

vBlades project

- The Itanium processor supports 4 privilege rings - PL0, PL1, PL2, PL3.
- Privileged instructions can only be executed by the kernel at PL0.
- App's run at PL3 & can execute only non-privileged instructions.
- The VMM uses ring compression & runs itself at PL0 & PL1 while forcing a guest OS to run at PL2.
- A problem - privilege leaking - several nonprivileged instructions allow an appⁿ to determine the current privilege level (CPL).
- Itanium was selected because of its multiple functional units & multithreading support.

CPU virtualization:

when a guest OS attempts to execute a privileged instruction the VMM traps & emulates, the instruction.

• **Complication** - Itanium doesn't have an *IR & the VMM has to use state info to determine whether an instr. is privileged.

• **Complication** - caused by *RSE, which operates concurrently w/ the processor & may attempt to access memory & generate a page fault.

- A no. of privileged-sensitive instr. behave differently as a function of the privilege level.

The VMM replaces each one of them w/ a privileged instr. during the dynamic transformation of the instruction stream.

contd.

contd.

- Memory virtualization is guided by the realization that a VMM should not be involved in most read & write operations to prevent a significant degradation of performance, but at the same time the VMM should exercise tight control & prevent a guest OS from acting maliciously.
- The vBlades VMM doesn't allow a guest OS to access memory directly. It inserts an additional layer of indirection called metaphysical addressing b/w virtual & real addressing.
- A guest OS is placed in metaphysical addressing mode. If the address is virtual, the VMM first checks if the guest OS is allowed to access that address, & if it is, it provides the regular address translation. If the address is physical, the VMM is not involved.

